US_ID	User Story	Scenario	Given	When	Then
US_01	As a student, I want to describe my academic situation in plain language, so that I can get a recommended schedule without filling out a rigid form.	Student types their situation in plain language	A student has written a free-text description of their program, completed courses, and constraints	They submit it to the extraction service	The system returns a structured object containing their program, level, completed courses, and credit limit
US_01		Course codes are written inconsistently	A student writes course codes in different formats, such as "CS 101" and "CS201"	The input is processed	All course codes are returned in a single consistent format
US_01		Output is handed to the solver	The extraction has completed	The result is returned	The object validates against the shared solver schema with no invalid fields
US_02	As a student, I want to optionally upload my transcript, so that I don't have to retype every course I've already completed.	Student uploads a transcript alongside their message	A student has a transcript listing courses they have completed	They submit it with their free-text message	The courses are extracted and added to the plan as suggestions
US_02		Transcript courses are not treated as final	Courses have been read from an uploaded transcript	The result is returned	Those courses are flagged for the student to confirm rather than accepted automatically
US_03	As an advisor, I want the system to flag anything the AI assumed or was unsure about, so that the student can confirm it instead of trusting a hidden guess.	The student leaves a required detail unstated	A student's message does not mention their starting term	The input is processed	The system fills a sensible default and records it in the assumptions list
US_03		Uncertain fields are surfaced for review	The system has filled in one or more values the student did not state	The result is returned	Those field names appear under needs_confirmation so the student can verify them
US_03		Input cannot be mapped to any field	A student includes information that does not fit the schema	The input is processed	That text is preserved in unparsed_notes rather than being discarded
US_04	As a student, I want to state my top priorities, so that the recommended schedule reflects what matters to me.	Student states what they care about most	A student mentions interests or priorities such as focusing on machine learning	The input is processed	Up to three priorities are captured in the order the student expressed them
US_04		Priorities reach the solver as soft preferences	Priorities have been captured	The object is passed to the solver	They influence how valid schedules are ranked, without affecting whether a schedule is valid
US_05	As a user, I want to flag incorrect or hallucinated AI output, so that the team can identify and improve weak spots.	User sees output that looks wrong	A result has been returned with assumption and confirmation flags	The user reviews it and finds an error	They can submit a report tied to that specific result for the team to review