

US_ID	Feature	User Story	Acceptance Criteria	Story Points
US_01	AI Input Extraction	As a student, I want to describe my academic situation in plain language, so that I can get a recommended schedule without filling out a rigid form.	- A free-text message returns a structured object with program, level, completed courses, and credit limit. - Loosely written codes ("CS 101") are normalized ("CS101"). - The output validates against the shared solver schema.	8
US_02	Transcript Ingestion	As a student, I want to optionally upload my transcript, so that I don't have to retype every course I've already completed.	- The endpoint accepts optional transcript text alongside the message. - Transcript courses are added as suggestions flagged for confirmation, not treated as authoritative. - PDF text is extracted before processing.	5
US_03	Transparency / Confidence Flags	As an advisor, I want the system to flag anything the AI assumed or was unsure about, so that the student can confirm it instead of trusting a hidden guess.	- Any default the AI fills in appears in an "assumptions" list. - Fields needing verification appear in "needs_confirmation". - Unmappable input is preserved in "unparsed_notes" rather than dropped.	5
US_04	Preference Weighting	As a student, I want to state my top priorities (e.g. "focus on ML"), so that the recommended schedule reflects what matters to me.	- Up to 3 priorities are captured, in order of importance. - Priorities are passed to the solver as soft preferences (affect ranking, not validity).	3
US_05	Feedback Mechanism	As a user, I want to flag incorrect or hallucinated AI output, so that the team can identify and improve weak spots.	- Each result surfaces the assumption/confirmation flags a feedback prompt can use. - A user can submit a flag tied to a specific result.	3
Total Story Points				24